

# Lecture 5: Floating Point Math

ME/AE 6705

Introduction to Mechatronics

Dr. Jonathan Rogers



# Lesson Objectives

- Understand how decimal numbers (floating point numbers) are represented on microprocessors
- Be able to convert a decimal number into floating point representation
- Understand the limits of precision using floating point numbers
- Understand fixed point representations
- Understand methods of explicit type conversion (casting) in C programming



# Limits of Integer Math

- Last class we learned how integer arithmetic is performed on microprocessors
- In many applications, integer math is not sufficient for the calculations we need

$$\frac{1024}{1099} = 0.9318$$


← Zero is often not an acceptable result

- Example: Robot determines its speed = 4 ft/s and needs to travel 3 ft. Does it have 0 sec or 0.75 sec left to travel?



# Floating Point Numbers


- Floating point numbers are a way to represent decimal numbers with binary digits
- Floating point numbers almost always use 32-bit or 64-bit representation
  - 64-bit can represent larger range of decimal numbers to higher precision
- IEEE 754 floating point standard (32-bit):



# Floating Point Numbers

- Decimal number  $A$  represented in floating point as follows:

$$A = (\text{Sign}) \times 2^{\text{exponent}} \times \text{mantissa}$$



# Floating Point Numbers

- Sign bit – 1 if negative number, 0 if positive number
- Mantissa – represents number between 1 and 2
  - Mantissa computed by letting each bit represent negative power of 2 in decreasing order

$$\begin{aligned}\text{Mantissa} &= 1 + (\text{bit } 22) \times 2^{-1} + (\text{bit } 21) \times 2^{-2} + \dots (\text{bit } 0) \times 2^{-23} \\ &= 1 + (\text{bit } 22) \times \frac{1}{2} + (\text{bit } 21) \times \frac{1}{4} + \dots (\text{bit } 0) \times \frac{1}{8388608} \\ &= 1 + \sum_{i=0}^{22} (\text{bit } i) \times 2^{i-23}\end{aligned}$$



# Floating Point Numbers

- Exponent represented by bits 23-30 (8 bits total)
  - Exponent is binary integer represented by 8 bits minus 127
  - The -127 is used so both positive and negative exponents can be represented
  - Value of exponent runs between 128 and -127

$$\text{Exponent} = \left( \sum_{i=23}^{30} (\text{bit } i) \times 2^{i-23} \right) - 127$$

$$A = (\text{Sign}) \times 2^{\text{Exponent}} \times \text{mantissa}$$



*Used here in calculation of floating point number*



# Floating Point Number Example

- Determine decimal value of following floating point number:

0 1 0 0 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

- Step 1: Number is positive since sign bit is 0.
- Step 2: Compute exponent value
  - Exponent =  $2^7 - 127 = 128 - 127 = 1$
- Step 3: Compute mantissa
  - Mantissa =  $1 + 1/2 + 1/4 + 1/8 + 1/16 = 1.9375$





# Floating Point Number Example

- Putting it all together:

$$\begin{aligned} A &= (\text{Sign}) \times 2^{\text{Exponent}} \times \text{mantissa} \\ &= 1 \times 2^1 \times 1.9375 = \mathbf{3.875} \end{aligned}$$

- Notes:
  - Range of numbers that 32-bit floating point can represent is  $\pm 1.18 \times 10^{-38}$  to  $\pm 3.4 \times 10^{38}$
  - Zero is represented by mantissa = 0 and exponent = 0 (special case)

# Floating Point Number Example

- Example problem: Determine the decimal number represented by:

1    0 1 0 1 0 0 1 1    0 0 0 0 1 1 0 1 0 0 1 0 1 0 0 0 0 0 0 1 1 0 0



# Floating Point Numbers

- Converting decimal number  $A$  to floating point representation is a bit trickier
  - Step 1: Note whether number is  $+$  or  $-$  to determine sign bit
  - Step 2: Let  $2^z = |A|$  and then determine exponent  $e$  as  $e = \text{floor}(z)$
  - Step 3: Determine mantissa by recursively adding negative powers of 2



# Floating Point Number Example

- Example problem: Convert 0.15625 into floating point representation.
  - Step 1: Sign bit will be 0 since number is positive
  - Step 2: Find the exponent

$$2^z = 0.15625$$

$$z = \log_2(0.15625) = \frac{\ln(0.15625)}{\ln(2)}$$
$$= -2.6780719$$

➔  $e = \text{floor}(z) = -3$



# Floating Point Number Example

- Example continued...
  - Step 3: Computing the mantissa

$$0.15625 = (1) (2^{-3}) x \quad \longleftarrow \text{(definition of floating point number, } x \text{ is mantissa)}$$

$$x = \frac{0.15625}{2^{-3}} = 1.25 \quad (x \text{ will always be } \geq 1)$$

- *Start at  $k=0$ , compare  $x-1$  to  $2^{-k}$ , if  $x \geq 2^{-k}$  then set  $k^{th}$  bit to 1, otherwise set it to 0*
- *Then set  $x = x - 2^{-k}$  and repeat until  $k = 23$  or  $x = 0$*

$$k = 0 \quad 1.25 - 1 \geq 2^{-1} \quad \text{no} \rightarrow \text{set } k^{th} \text{ bit to 0} \quad x = 1.25 - 1 = 0.25$$

$$k = 1 \quad 0.25 \geq 2^{-2} \quad \text{yes} \rightarrow \text{set } k^{th} \text{ bit to 1} \quad x = 0.25 - 0.25 = 0$$

# Floating Point Number Example

- Example continued...
- Putting everything together, 1.25 is represented in 32-bit floating point format as

0 0 1 1 1 1 1 0 0 0 1 0



Note: This is  $-3+127 = 124$

- Writing this in hexadecimal,

0 0 1 1 1 1 1 0 0 0 1 0



# Floating Point Number Example

- Example continued...
- Putting everything together, 1.25 is represented in 32-bit floating point format as

0 0 1 1 1 1 1 0 0 0 1 0

- Writing this in hexadecimal,

0 0 1 1 1 1 1 0 0 0 1 0

3E 20 00 00



# Floating Point Precision

- 32-bit floating point format can only represent a number exactly if it can be created by summing powers of 2 from  $2^{-1}$  to  $2^{-23}$ 
  - i.e., using 24 bits
  - In previous example, our number 0.15625 could be represented exactly because  $0.15625 = 2^{-3} + 2^{-5}$
- If the above is not true for the number we are trying to represent, the number will be approximated in floating point by truncating it at the 24<sup>th</sup> bit
  - This means it will be correct to the ~7<sup>th</sup> decimal place





# Precision Example

- Example: How precisely is  $16\pi$  represented in floating point format (32 bit)?
- Solution:
  - By repeating steps from previous example, we find that the floating point representation of  $16\pi$  is

sign bit + exponent

0x84490FDB

mantissa



# Precision Example

- The exponent value is 0x84 which is 132 in decimal, thus

$$\text{exponent} = 132 - 127 = 5$$

- The mantissa (or fraction) is 0x490FDB which in binary is

1 0 0   1 0 0 1   0 0 0 0   1 1 1 1   1 1 0 1   1 0 1 1

4        9        0        F        D        B



# Precision Example


- Writing the mantissa  $f$  in decimal,

$$\begin{aligned} f &= 1 + 2^{-1} + 2^{-4} + 2^{-7} + 2^{-12} + 2^{-13} + 2^{-14} + \\ &\quad 2^{-15} + 2^{-16} + 2^{-17} + 2^{-19} + 2^{-20} + 2^{-22} + 2^{-23} \\ &= 1.5707963705 \end{aligned}$$

- Thus, the final decimal value of the floating point representation is:

$$1 \times 2^5 \times 1.5707963705 = 50.265483856$$


$$\text{Actual value of } 16\pi = 50.265482457$$



Correct to ~6<sup>th</sup>  
decimal place

# Floating Point Precision

- So far, we have discussed 32-bit floating point numbers
  - In C, these are numbers of type “float”
- In addition, 64-bit floating point numbers are defined to give higher precision
  - In C, these are numbers of type “double”
  - Uses 11-bit exponent instead of 8
  - Uses 52-bit mantissa (or fraction) instead of 23



# Floating Point Precision

- Double precision (64-bit) floating point allows much greater precision and range compared to 32-bit

Type	Width	Range at Full Precision	Precision
Single precision (float)	32 bits	$\pm 10^{-38}$ to $\pm 10^{38}$	Approximately 7 decimal digits
Double precision (double)	64 bits	$\pm 10^{-305}$ to $\pm 10^{305}$	Approximately 15 decimal digits



# Fixed Point Numbers

- Processing floating point numbers (adding, multiplying, dividing, etc) requires a special set of registers on the microprocessor
  - This is called a Floating Point Unit (FPU)
  - Many microcontrollers do not have FPU's
  - If FPU is not available or not activated, processor cannot handle floating point numbers
- Furthermore, floating point math is slower than integer math even if FPU is available
  - Thus, alternative called Fixed Point sometimes used



# Fixed Point Numbers

- Fixed point numbers are used to represent non-integer values
- They consist of an integer  $I$  and a fixed constant  $\Delta$ 
  - Integer part  $I$  is stored in memory as normal integer type
  - Constant  $\Delta$  is not stored in memory and cannot be changed, but rather is set by programmer when software is written

$$f = I \times \Delta$$



# Fixed Point Numbers

- $\Delta$  is a multiplier value that we use during input and output to convert physical units to integers which are stored in memory and manipulated
  - Example:  $\Delta = 0.001$ . Thus, to represent 2.314 in memory, we would simply use an integer value of 2314.
  - Before inputting values to program, we would divide our values by 0.001.
  - After outputting values from program, we multiply integer values by 0.001.





# Fixed Point Numbers

- Decimal fixed point number ( $\Delta$  power of 10)
  - Easier for input and output

Decimal fixed point number =  $I \times 10^m$  for some fixed constant  $m < 0$

- Binary fixed point number ( $\Delta$  power of 2)
  - Easier for mathematical calculations

Binary fixed point number =  $I \times 2^n$  for some fixed constant  $n < 0$

- Note: Binary fixed point numbers are easier to use if you have different values of  $\Delta$  throughout your program and you need to convert between them.

# Fixed Point Example

- Example: Suppose for your program you decide to use  $\Delta = 0.001$ . How would  $3\pi$  be represented in fixed point format in your program?
  - Furthermore, what size integer is required to represent it?



# Fixed Point Example

- Example: Suppose for your program you decide to use  $\Delta = 0.001$ . How would  $3\pi$  be represented in fixed point format in your program?
  - Furthermore, what size integer is required to represent it?
- Solution:

$$3\pi = 9.42477796076938$$

$$\rightarrow 9.42477796076938 / 0.001 = 9424.7779607...$$

- Thus, an integer of 9425 would be used to represent  $3\pi$  in memory
- Note this integer must be at least 16 bits to avoid problems with overflow

# Fixed Point Example

- Example: Suppose for your program you decide to use  $\Delta = 2^{-8}$ . How would  $3\pi$  be represented in fixed point format in your program?
  - Furthermore, what size integer is required to represent it?



# Fixed Point Example

- Example: Suppose for your program you decide to use  $\Delta = 2^{-8}$ . How would  $3\pi$  be represented in fixed point format in your program?
  - Furthermore, what size integer is required to represent it?

- Solution:

$$\begin{aligned} 3\pi / 2^{-8} &= 9.4247779607 / 0.00390625 \\ &= 2412.74315795... \end{aligned}$$

- Thus, an integer of value 2413 would be used to represent  $3\pi$  in memory
- Note this integer must be at least 16 bits to avoid problems with overflow

# Fixed Point Arithmetic

- $\Delta$  can be thought of as just a “scale factor” that converts units in your program to physical units
- For example, you may be dealing with Volts in the “physical” domain, but processor will deal with millivolts if  $\Delta = 0.001$ .
- As a result, fixed point numbers can be added, subtracted, multiplied, and divided exactly the same as with integers
  - If they have different values of  $\Delta$ , conversion to same  $\Delta$  must be done first



# Fixed Point Example

- Suppose you have a value of  $B = 4386$  in memory, which is a fixed point value with  $\Delta = 0.001$  representing voltage measured across two pins
- You have another value of  $C = 357421$  which is fixed point value with  $\Delta = 10^{-6}$  representing voltage measured across two other pins
- You wish to add these voltages in your program so that  $D = B + C$  is a fixed point value. How do you do this so as to maintain maximum precision (minimum truncation error)?



# Fixed vs Floating Point

## Floating Point

- Uses more memory and is slower
  - All numbers use at least 32 bits
- Must have FPU available and active
- Can handle large variation in range of numbers
- Called “floating point” because number of decimal places varies depending on how large number is


## Fixed Point

- Uses less memory and is faster
- Will always work regardless of processor (even if no FPU available)
- Only accurate if range of numbers is known and small
- Called “fixed point” because number of decimal places always fixed for each number



# Type Conversions

- C99 compilers will automatically perform type conversions when arithmetic is done on two different variable types
- When two types are operated on, lower order is converted to higher with result stored as higher order
- Order is:
  - Double
  - Float
  - Signed 32 bit
  - Unsigned 32 bit
  - Signed 16 bit
  - Unsigned 16 bit
  - Signed 8 bit
  - Unsigned 8 bit



# Type Conversions

- Consider fragment of C program below:

```
int8_t num1 = 77 ;  
int8_t num2 = 82 ;  
int16_t num3 = -534 ;  
int32_t res = 0 ;
```

```
res = num3 - (num1 + num2) ;
```

Added as 8 bit, stored as 8 bit

8 bit converted to 16 bit,  
added as 16 bit

16 bit converted to 32 bit



# Type Conversions

- Explicit conversion from one variable to type to another can be done at any time by placing (type) directly before variable
  - This is called typecasting
- Consider previous case. Adding num1 and num2 would cause overflow (since  $77 + 82 > 127$ )
- Can fix this by typecasting one number to 16-bit integer

```
int8_t num1 = 77 ;  
int8_t num2 = 82 ;  
int16_t num3 = -534 ;  
int32_t res = 0 ;  
  
res = (int32_t)num3 - ((int16_t)num1 + num2) ;
```

# Type Conversions

- Similarly, typecasting is also done for 32 bit and 64 bit floating point types

```
float num1 = 77.6598 ;  
float num2 = 82.4758 ;  
int16_t num3 = -534 ;  
double res = 0 ;
```

```
res = num3 - (num1 / num2) ;
```

- num1 and num2 divided as 32 bit floats, num3 converted to 32 bit float and added to  $-(\text{num1} + \text{num2})$
- Result then converted to 64 bit float and set equal to res



# Type Conversions

- Using explicit typecasting is always better as it helps to keep track of precision during math operations

```
float num1 = 77.6598 ;  
float num2 = 82.4758 ;  
int16_t num3 = -534 ;  
double res = 0 ;
```

```
res = (double) ((float) num3 - (num1 / num2)) ;
```



# printf Function in C

- In C, you can print variables to screen during runtime using **printf** function
  - Supported for TI microcontroller in debug mode
- Example functionality:

```
printf("Value of num1 is %u\n", num1);
```

← Prints num1 to screen,  
then ends line (return).

```
printf("num2 = %f   num3 = %.10f\n", num2, num3);
```

Important: %XX characters tell compiler how to interpret value during printing process (i.e., variable type).

← Prints num2 and  
num3 to screen, then  
ends line (return).

# printf Function in C

Specifier	Output	Example
%c	Character	a
%d or %i	Signed decimal integer	-392
%ld	Signed 32 bit long decimal integer	65846214
%e	Scientific notation	6.022141e23
%E	Scientific notation, capital letter	6.022141E23
%f	Floating point	3.14159
%o	Unsigned octal	610
%s	String of characters	Sample
%u	Unsigned decimal integer	7235
%x	Unsigned hexadecimal integer	7fa
%X	Unsigned hexadecimal integer (capital letters)	7FA
%4.10f	Specific precision floating point	0014.6589743241



# printf Function in C

- Note: Printing out using the wrong specifier for the type of variable can produce some very strange results

```
int32_t res = 15 ;  
printf("res = %f\n", res) ;
```

→ prints res = 0.000

```
int32_t res = 15 ;  
printf("res = %d\n", res) ;
```

→ prints res = 15





# Review Example Code

- Review Lecture 5 example code

